

# Dataproc: Qwik Start - Command Line

GSP104



- Dataproc provides a fully managed cloud service for running Apache Spark and Hadoop clusters quickly and cost-efficiently.
- Operations that once took hours or days now complete in seconds or minutes with automatic cluster scaling.
- This lab demonstrates creating a Dataproc cluster entirely through command line interface.
- You'll execute a simple Apache Spark job directly within the provisioned cluster environment.
- Finally, resize the cluster by modifying the number of worker nodes dynamically.



## Dataproc: Qwik Start - Command Line

00:37:18  
End Lab  
Time limit

gcloud is the command-line tool for Google Cloud. It comes pre-installed on Cloud Shell and supports tab-completion.

0/100

3. (Optional) You can list the active account name with this command:

```
gcloud auth list
```

4. Click **Authorize**.

**Output:**

```
ACTIVE: *
ACCOUNT: student-02-77569a450c62@qwiklabs.net

To set the active account, run:
$ gcloud config set account `ACCOUNT`
```

5. (Optional) You can list the project ID with this command:

```
gcloud config list project
```

Lab instructions and tasks

GSP104

Overview

Setup and requirements

Task 1. Create a cluster

Task 2. Submit a job

Task 3. Update a cluster

Task 4. Test your understanding

Congratulations!

Caution: When you are in the console, do not deviate from the lab instructions. Doing so may cause your account to be blocked.  
[Learn more.](#)

[Open Google Cloud console](#)

Username

student-02-77569a450c62@



Password

1gd108qixhdc



Project ID

qwiklabs-gcp-03-9a34201



## Student Resources

[Dataproc: Qwik Start - Qwiklabs](#)[Previous](#)[Next](#)

Dashboard – qwiklabs-gcp-03-9... X +

https://console.cloud.google.com/home/dashboard?project=qwiklabs-gcp-03-9a342013a3da&pli=1&cloudshell=true

Free AI Paraphrasin... Transcript of Case St... GCP-LAB Hydra datawarehouse Products

Google Cloud qwiklabs-gcp-03-9a342013a3da Search (/) for resources, docs, products, and more Search

CLOUD SHELL Terminal (qwiklabs-gcp-03-9a342013a3da) X + Open Editor

Gemini CLI is available in Cloud Shell terminal! Type gemini to try it. [Learn more](#) Don't show again Dismiss

Welcome to Cloud Shell! Type "help" to get started, or type "gemini" to try prompting with Gemini CLI.  
Your Cloud Platform project in this session is set to **qwiklabs-gcp-03-9a342013a3da**.  
Use `gcloud config set project [PROJECT\_ID]` to change to a different project.  
student\_02\_77569a450c62@cloudshell:~ (qwiklabs-gcp-03-9a342013a3da)\$ gcloud auth list  
Credentials  
ACTIVE: \*  
ACCOUNT: student-02-77569a450c62@qwiklabs.net  
  
To set the active account, run:  
\$ gcloud config set account `ACCOUNT`  
  
student\_02\_77569a450c62@cloudshell:~ (qwiklabs-gcp-03-9a342013a3da)\$ gcloud config list project  
[core]  
project = qwiklabs-gcp-03-9a342013a3da  
  
Your active configuration is: [cloudshell-22103]  
student\_02\_77569a450c62@cloudshell:~ (qwiklabs-gcp-03-9a342013a3da)\$

News for you HC erred with na... Search

04:05 PM 28-12-2025 ENG US

00:36:58  
End Lab  
Time limit

Caution: When you are in the console, do not deviate from the lab instructions. Doing so may cause your account to be blocked.

[Learn more.](#)[Open Google Cloud console](#)

Username

student-02-77569a450c62



Password

1gd108qixhdc



Project ID

qwiklabs-gcp-03-9a34201



Student Resources

[Dataproc: Qwik Start - Qwiklabs](#)

## Task 1. Create a cluster

1. In Cloud Shell, run the following command to set the Region:

```
gcloud config set dataproc/region us-east1
```

2. Disable the Dataproc API:

```
gcloud services disable dataproc.googleapis.com --force
```

3. Re-enable the Dataproc API:

```
gcloud services enable dataproc.googleapis.com
```

4. Dataproc creates staging and temp buckets that are shared among clusters in the same region. Since we're not specifying an account for Dataproc to use, it will use the Compute Engine default service account, which doesn't have storage bucket permissions by default. Let's add those.

Lab instructions and tasks

0/100

GSP104

Overview

Setup and requirements

Task 1. Create a cluster

Task 2. Submit a job

Task 3. Update a cluster

Task 4. Test your understanding

Congratulations!

[Previous](#)[Next](#)

Dashboard – qwiklabs-gcp-03-9... X +

https://console.cloud.google.com/home/dashboard?project=qwiklabs-gcp-03-9a342013a3da&pli=1&cloudshell=true

Free AI Paraphrasin... Transcript of Case St... GCP-LAB Hydra datawarehouse Products

Google Cloud qwiklabs-gcp-03-9a342013a3da Search (/) for resources, docs, products, and more Search

CLOUD SHELL Terminal (qwiklabs-gcp-03-9a342013a3da) X + Open Editor

Gemini CLI is available in Cloud Shell terminal! Type gemini to try it. [Learn more](#)

Don't show again Dismiss

```
student_02_77569a450c62@cloudshell:~ (qwiklabs-gcp-03-9a342013a3da)$ gcloud config set dataproc/region us-east1
Updated property [dataproc/region].
student_02_77569a450c62@cloudshell:~ (qwiklabs-gcp-03-9a342013a3da)$ gcloud services disable dataproc.googleapis.com --force
Operation "operations/acat.p17-1019275457196-1acb3562-08be-462e-820c-896cedff2ald" finished successfully.
student_02_77569a450c62@cloudshell:~ (qwiklabs-gcp-03-9a342013a3da)$ gcloud services enable dataproc.googleapis.com
Operation "operations/acf.p2-1019275457196-0f855cad-2a40-453c-b435-3e9cc08e6ae0" finished successfully.
student_02_77569a450c62@cloudshell:~ (qwiklabs-gcp-03-9a342013a3da)$ █
```

Top Stories 4 Did Osman Hadi'...

Search

04:06 PM 28-12-2025 ENG US

00:36:50  
End Lab  
Time limit

Caution: When you are in the console, do not deviate from the lab instructions. Doing so may cause your account to be blocked.

[Learn more.](#)[Open Google Cloud console](#)

Username

student-02-77569a450c62



Password

1gd108qixhdc



Project ID

qwiklabs-gcp-03-9a34201



Student Resources

[Dataproc: Qwik Start - Qwiklabs](#)

4. Dataproc creates staging and temp buckets that are shared among clusters in the same region. Since we're not specifying an account for Dataproc to use, it will use the Compute Engine default service account, which doesn't have storage bucket permissions by default. Let's add those.

- First, run the following commands to grab the PROJECT\_ID and PROJECT\_NUMBER:

```
PROJECT_ID=$(gcloud config get-value project) && \
gcloud config set project $PROJECT_ID
```

```
PROJECT_NUMBER=$(gcloud projects describe $PROJECT_ID -- \
format='value(projectNumber)')
```

- Now run the following command to give the Storage Admin and Dataproc Worker role to the Compute Engine default service account:

```
gcloud projects add-iam-policy-binding $PROJECT_ID \
--member=serviceAccount:$PROJECT_NUMBER-
compute@developer.gserviceaccount.com \
--role=roles/storage.admin
```

```
gcloud projects add-iam-policy-binding $PROJECT_ID \
```

Lab instructions and tasks

0/100

GSP104

Overview

Setup and requirements

Task 1. Create a cluster

Task 2. Submit a job

Task 3. Update a cluster

Task 4. Test your understanding

Congratulations!

[Previous](#)[Next](#)

Dashboard - qwiklabs-gcp-03-9... X +

https://console.cloud.google.com/home/dashboard?project=qwiklabs-gcp-03-9a342013a3da&pli=1&cloudshell=true

Free AI Paraphrasin... Transcript of Case St... GCP-LAB Hydra datawarehouse Products

Google Cloud qwiklabs-gcp-03-9a342013a3da Search (/) for resources, docs, products, and more Search

CLOUD SHELL Terminal (qwiklabs-gcp-03-9a342013a3da) X + Open Editor

Gemini CLI is available in Cloud Shell terminal! Type gemini to try it. [Learn more](#)

Don't show again Dismiss

```
student_02_77569a450c62@cloudshell:~ (qwiklabs-gcp-03-9a342013a3da)$ PROJECT_ID=$(gcloud config get-value project) && \
gcloud config set project $PROJECT_ID

PROJECT_NUMBER=$(gcloud projects describe $PROJECT_ID --format='value(projectNumber)')
Your active configuration is: [cloudshell-22103]
Updated property [core/project].
student_02_77569a450c62@cloudshell:~ (qwiklabs-gcp-03-9a342013a3da)$
student_02_77569a450c62@cloudshell:~ (qwiklabs-gcp-03-9a342013a3da)$ gcloud projects add-iam-policy-binding $PROJECT_ID \
--member=serviceAccount:$PROJECT_NUMBER-compute@developer.gserviceaccount.com \
--role=roles/storage.admin

gcloud projects add-iam-policy-binding $PROJECT_ID \
--member=serviceAccount:$PROJECT_NUMBER-compute@developer.gserviceaccount.com \
--role=roles/dataproc.worker
Updated IAM policy for project [qwiklabs-gcp-03-9a342013a3da].
bindings:
- members:
  - serviceAccount:qwiklabs-gcp-03-9a342013a3da@qwiklabs-gcp-03-9a342013a3da.iam.gserviceaccount.com
    role: roles/bigquery.admin
- members:
  - serviceAccount:1019275457196@cloudbuild.gserviceaccount.com
    role: roles/cloudbuild.builds.builder
- members:
  - serviceAccount:service-1019275457196@gcp-sa-cloudbuild.iam.gserviceaccount.com
    role: roles/cloudbuild.serviceAgent
- members:
  - serviceAccount:service-1019275457196@compute-system.iam.gserviceaccount.com
    role: roles/compute.serviceAgent
- members:
```

5 17°C Sunny

Search

04:07 PM 28-12-2025 ENG US

Dashboard – qwiklabs-gcp-03-9... X +

https://console.cloud.google.com/home/dashboard?project=qwiklabs-gcp-03-9a342013a3da&pli=1&cloudshell=true

Free AI Paraphrasing... Transcript of Case St... GCP-LAB Hydra datawarehouse Products

Google Cloud qwiklabs-gcp-03-9a342013a3da Search (/) for resources, docs, products, and more Search

CLOUD SHELL Terminal (qwiklabs-gcp-03-9a342013a3da) + Open Editor

i Gemini CLI is available in Cloud Shell terminal! Type gemini to try it. [Learn more](#)

Don't show again Dismiss

```
role: roles/cloudbuild.serviceAgent
- members:
  - serviceAccount:service-1019275457196@compute-system.iam.gserviceaccount.com
  role: roles/compute.serviceAgent
- members:
  - serviceAccount:service-1019275457196@dataproc-accounts.iam.gserviceaccount.com
  role: roles/dataproc.serviceAgent
- members:
  - serviceAccount:1019275457196-compute@developer.gserviceaccount.com
  role: roles/dataproc.worker
- members:
  - serviceAccount:1019275457196-compute@developer.gserviceaccount.com
  - serviceAccount:1019275457196@cloudservices.gserviceaccount.com
  role: roles/editor
- members:
  - serviceAccount:admiral@qwiklabs-services-prod.iam.gserviceaccount.com
  - serviceAccount:qwiklabs-gcp-03-9a342013a3da@qwiklabs-gcp-03-9a342013a3da.iam.gserviceaccount.com
  - user:student-02-77569a450c62@qwiklabs.net
  role: roles/owner
- members:
  - serviceAccount:1019275457196-compute@developer.gserviceaccount.com
  - serviceAccount:qwiklabs-gcp-03-9a342013a3da@qwiklabs-gcp-03-9a342013a3da.iam.gserviceaccount.com
  role: roles/storage.admin
- members:
  - user:student-02-77569a450c62@qwiklabs.net
  role: roles/viewer
etag: BwZHALeX-00=
version: 1
student_02_77569a450c62@cloudshell:~ (qwiklabs-gcp-03-9a342013a3da)$
```

5 17°C Sunny

Search

04:07 PM 28-12-2025



## Dataproc: Qwik Start - Command Line

00:36:43

Time limit

Caution: When you are in the console, do not deviate from the lab instructions. Doing so may cause your account to be blocked.

[Learn more.](#)[Open Google Cloud console](#)

Username

student-02-77569a450c62



Password

1gd108qixhdc



Project ID

qwiklabs-gcp-03-9a34201



## Student Resources

[Dataproc: Qwik Start - Qwiklabs](#)

```
PROJECT_NUMBER=$(gcloud projects describe $PROJECT_ID --format='value(projectNumber)')
```

- Now run the following command to give the Storage Admin and Dataproc Worker role to the Compute Engine default service account:

```
gcloud projects add-iam-policy-binding $PROJECT_ID \
--member=serviceAccount:$PROJECT_NUMBER-
compute@developer.gserviceaccount.com \
--role=roles/storage.admin
```

```
gcloud projects add-iam-policy-binding $PROJECT_ID \
--member=serviceAccount:$PROJECT_NUMBER-
compute@developer.gserviceaccount.com \
--role=roles/dataproc.worker
```

- Enable Private Google Access on your subnetwork by running the following command:

```
gcloud compute networks subnets update default --region=us-east1
--enable-private-ip-google-access
```

Lab instructions and tasks

0/100

GSP104

Overview

Setup and requirements

Task 1. Create a cluster

Task 2. Submit a job

Task 3. Update a cluster

Task 4. Test your understanding

Congratulations!

[Previous](#)[Next](#)



## Dataproc: Qwik Start - Command Line

00:36:38

Time limit

Caution: When you are in the console, do not deviate from the lab instructions. Doing so may cause your account to be blocked.

[Learn more.](#)[Open Google Cloud console](#)

Username

student-02-77569a450c62



Password

1gd108qixhdc



Project ID

qwiklabs-gcp-03-9a34201



## Student Resources

[Dataproc: Qwik Start - Qwiklabs](#)[Previous](#)

0/100

Lab instructions and tasks

```
compute@developer.gserviceaccount.com \
--role=roles/dataproc.worker
```

3. Enable Private Google Access on your subnetwork by running the following command:

```
gcloud compute networks subnets update default --region=us-east1 \
--enable-private-ip-google-access
```

4. Run the following command to create a cluster called `example-cluster` with `e2-standard-4` VMs and default Cloud Dataproc settings:

```
gcloud dataproc clusters create example-cluster --worker-boot-disk-size 500 --worker-machine-type=e2-standard-4 --master-machine-type=e2-standard-4
```

5. If asked to confirm a zone for your cluster. Enter `Y`.

Your cluster will build for a couple of minutes.

[Next](#)

Dashboard – qwiklabs-gcp-03-9... X +

https://console.cloud.google.com/home/dashboard?project=qwiklabs-gcp-03-9a342013a3da&pli=1&cloudshell=true

Free AI Paraphrasin... Transcript of Case St... GCP-LAB Hydra datawarehouse Products

Google Cloud qwiklabs-gcp-03-9a342013a3da Search (/) for resources, docs, products, and more Search

CLOUD SHELL Terminal (qwiklabs-gcp-03-9a342013a3da) X + Open Editor

Gemini CLI is available in Cloud Shell terminal! Type gemini to try it. [Learn more](#) Don't show again Dismiss

```
student_02_77569a450c62@cloudshell:~ (qwiklabs-gcp-03-9a342013a3da)$ gcloud compute networks subnets update default --region=us-east1 --enable-private-ip-google-access
Updated [https://www.googleapis.com/compute/v1/projects/qwiklabs-gcp-03-9a342013a3da/regions/us-east1/subnetworks/default].
student_02_77569a450c62@cloudshell:~ (qwiklabs-gcp-03-9a342013a3da)$
student_02_77569a450c62@cloudshell:~ (qwiklabs-gcp-03-9a342013a3da)$ gcloud dataproc clusters create example-cluster --worker-boot-disk-size 500 --worker-machine-type=e2-standard-4 --master-machine-type=e2-standard-4
Waiting on operation [projects/qwiklabs-gcp-03-9a342013a3da/regions/us-east1/operations/090f61a0-af06-3a2b-a7cc-3651a0f9c504].
Waiting for cluster creation operation...
WARNING: For PD-Standard without local SSDs, we strongly recommend provisioning 1TB or larger to ensure consistently high I/O performance. See https://cloud.google.com/compute/docs/disks/performance for information on disk I/O performance.
WARNING: The firewall rules for specified network or subnetwork would allow ingress traffic from 0.0.0.0/0, which could be a security risk.
WARNING: No image specified. Using the default image version. It is recommended to select a specific image version in production, as the default image version may change at any time.
Waiting for cluster creation operation...done.
Created [https://dataproc.googleapis.com/v1/projects/qwiklabs-gcp-03-9a342013a3da/regions/us-east1/clusters/example-cluster] Cluster placed in zone [us-east1-c].
student_02_77569a450c62@cloudshell:~ (qwiklabs-gcp-03-9a342013a3da)$
```

5 17°C Sunny

Search

04:13 PM 28-12-2025



Free AI Paraphrasin...

Transcript of Case St...

GCP-LAB

Hydra

datawarehouse

Products

Google Skills

Partner



6365



Contents

Dataproc: Qwik Start - Command Line

Dashboard

Catalog

Paths

Collections

00:28:44  
End Lab

Time limit

Caution: When you are in the console, do not deviate from the lab instructions. Doing so may cause your account to be blocked.

[Learn more.](#)

[Open Google Cloud console](#)

Username

student-02-77569a450c62



Password

1gd108qixhdc



Project ID

qwiklabs-gcp-03-9a34201



Student Resources

[Dataproc: Qwik Start - Qwiklabs](#)[Previous](#)

50/100

Lab instructions and task

GSP104

Overview

Setup and requirements

Task 1. Create a cluster

Task 2. Submit a job

Task 3. Update a cluster

Task 4. Test your understanding

Congratulations!

## Task 2. Submit a job

- Run this command to submit a sample Spark job that calculates a rough value for pi:

```
gcloud dataproc jobs submit spark --cluster example-cluster \
    --class org.apache.spark.examples.SparkPi \
    --jars file:///usr/lib/spark/examples/jars/spark-examples.jar
    -- 1000
```

The command specifies:

- That you want to run a `spark` job on the `example-cluster` cluster
- The `class` containing the main method for the job's pi-calculating application
- The location of the jar file containing your job's code
- The parameters you want to pass to the job—in this case, the number of tasks, which is `1000`

**Note:** Parameters passed to the job must follow a double dash (`--`). See the [gcloud documentation](#) for more information.

The job's running and final output is displayed in the terminal window:

[Next](#)



Caution: When you are in the console, do not deviate from the lab instructions. Doing so may cause your account to be blocked.

[Learn more.](#)[Open Google Cloud console](#)

Username

student-02-77569a450c62



Password

1gd108qixhdc



Project ID

qwiklabs-gcp-03-9a34201



Student Resources

[Dataproc: Qwik Start - Qwiklabs](#)[Previous](#)

50/100

Lab instructions and task

GSP104

Overview

Setup and requirements

Task 1. Create a cluster

Task 2. Submit a job

Task 3. Update a cluster

Task 4. Test your understanding

Congratulations!

- The location of the jar file containing your job's code
- The parameters you want to pass to the job—in this case, the number of tasks, which is 1000

**Note:** Parameters passed to the job must follow a double dash (--) See the [gcloud documentation](#) for more information.

The job's running and final output is displayed in the terminal window:

```
Waiting for job output...
...
Pi is roughly 3.14118528
...
state: FINISHED
```

## Test completed task

Click **Check my progress** to verify your performed task. If you have successfully submitted a job, you will see an assessment score.

[Next](#)

Dashboard – qwiklabs-gcp-03-9... X +

https://console.cloud.google.com/home/dashboard?project=qwiklabs-gcp-03-9a342013a3da&pli=1&cloudshell=true

Free AI Paraphrasin... Transcript of Case St... GCP-LAB Hydra datawarehouse Products

Google Cloud qwiklabs-gcp-03-9a342013a3da Search (/) for resources, docs, products, and more Search

CLOUD SHELL Terminal (qwiklabs-gcp-03-9a342013a3da) + Open Editor

Gemini CLI is available in Cloud Shell terminal! Type gemini to try it. [Learn more](#)

Don't show again Dismiss

```
- serviceAccount:service-1019275457196@gcp-sa-cloudbuild.iam.gserviceaccount.com
role: roles/cloudbuild.serviceAgent
student_02_77569a450c62@cloudshell:~ (qwiklabs-gcp-03-9a342013a3da)$ gcloud dataproc jobs submit spark --cluster example-cluster \
--class org.apache.spark.examples.SparkPi \
--jars file:///usr/lib/spark/examples/jars/spark-examples.jar -- 1000
Job [7230837507b64798a2feaa535de063e7] submitted.
Waiting for job output...
25/12/28 10:45:01 INFO SparkEnv: Registering MapOutputTracker
25/12/28 10:45:01 INFO SparkEnv: Registering BlockManagerMaster
25/12/28 10:45:01 INFO SparkEnv: Registering BlockManagerMasterHeartbeat
25/12/28 10:45:01 INFO SparkEnv: Registering OutputCommitCoordinator
25/12/28 10:45:02 INFO MetricsConfig: Loaded properties from hadoop-metrics2.properties
25/12/28 10:45:02 INFO MetricsSystemImpl: Scheduled Metric snapshot period at 10 second(s).
25/12/28 10:45:02 INFO MetricsSystemImpl: google-hadoop-file-system metrics system started
25/12/28 10:45:02 INFO DataprocSparkPlugin: Registered 188 driver metrics
25/12/28 10:45:03 INFO DefaultNoHARMFailoverProxyProvider: Connecting to ResourceManager at example-cluster-m.us-east1-c.c.qwiklabs-gcp-03-9a342013a3da.int
ernal./10.142.0.4:8032
25/12/28 10:45:03 INFO AHSProxy: Connecting to Application History server at example-cluster-m.us-east1-c.c.qwiklabs-gcp-03-9a342013a3da.internal./10.142.0
.4:10200
25/12/28 10:45:04 INFO Configuration: resource-types.xml not found
25/12/28 10:45:04 INFO ResourceUtils: Unable to find 'resource-types.xml'.
25/12/28 10:45:06 INFO YarnClientImpl: Submitted application application_1766918486251_0001
25/12/28 10:45:07 INFO DefaultNoHARMFailoverProxyProvider: Connecting to ResourceManager at example-cluster-m.us-east1-c.c.qwiklabs-gcp-03-9a342013a3da.int
ernal./10.142.0.4:8030
25/12/28 10:45:09 INFO GoogleCloudStorageImpl: Ignoring exception of type GoogleJsonResponseException; verified object already exists with desired state.
25/12/28 10:45:10 INFO GoogleHadoopOutputStream: hflush(): No-op due to rate limit (RateLimiter[stableRate=0.2qps]): readers will *not* yet see flushed dat
a for gs://dataproc-temp-us-east1-1019275457196-8ertei0p/3ef89636-2f20-4011-ac18-2c8259ea8458/spark-job-history/application_1766918486251_0001.inprogress [C
ONTEXT ratelimit_period="1 MINUTES" ]
Pi is roughly 3.141666871416669
```

5 17°C Sunny

Search

04:15 PM 28-12-2025

Dashboard - qwiklabs-gcp-03-9... X +

https://console.cloud.google.com/home/dashboard?project=qwiklabs-gcp-03-9a342013a3da&pli=1&cloudshell=true

Free AI Paraphrasing... Transcript of Case St... GCP-LAB Hydra datawarehouse Products

Google Cloud qwiklabs-gcp-03-9a342013a3da Search (/) for resources, docs, products, and more Search

CLOUD SHELL Terminal (qwiklabs-gcp-03-9a342013a3da) + Open Editor

Gemini CLI is available in Cloud Shell terminal! Type gemini to try it. [Learn more](#)

Don't show again Dismiss

```
25/12/28 10:45:03 INFO AHSPProxy: Connecting to Application History server at example-cluster-m.us-east1-c.c.qwiklabs-gcp-03-9a342013a3da.internal./10.142.0.4:10200
25/12/28 10:45:04 INFO Configuration: resource-types.xml not found
25/12/28 10:45:04 INFO ResourceUtils: Unable to find 'resource-types.xml'.
25/12/28 10:45:06 INFO YarnClientImpl: Submitted application application_1766918486251_0001
25/12/28 10:45:07 INFO DefaultNoHARMFailoverProxyProvider: Connecting to ResourceManager at example-cluster-m.us-east1-c.c.qwiklabs-gcp-03-9a342013a3da.internal./10.142.0.4:8030
25/12/28 10:45:09 INFO GoogleCloudStorageImpl: Ignoring exception of type GoogleJsonResponseException; verified object already exists with desired state.
25/12/28 10:45:10 INFO GoogleHadoopOutputStream: hflush(): No-op due to rate limit (RateLimiter[stableRate=0.2qps]): readers will *not* yet see flushed data for gs://dataproc-temp-us-east1-1019275457196-8ertei0p/3ef89636-2f20-4011-ac18-2c8259ea8458/spark-job-history/application_1766918486251_0001.inprogress [CONTEXT ratelimit_period="1 MINUTES"]
Pi is roughly 3.141666871416669
25/12/28 10:45:29 INFO DataprocSparkPlugin: Shutting down driver plugin. metrics=[action_http_patch_request=0, files_created=1, gcs_api_server_timeout_count=0, op_get_list_status_result_size=0, op_open=0, action_http_delete_request=2, gcs_api_time=1736, gcs_backoff_count=0, gcs_api_client_unauthorized_response_count=0, stream_read_close_operations=0, stream_read_bytes_backwards_on_seek=0, gs_filesystem_create=3, exception_count=0, gcs_exception_count=0, gcs_api_total_request_count=33, op_create=1, stream_read_vectored_operations=0, gcs_metadata_request=17, gcs_api_client_bad_request_count=0, action_http_put_request=3, op_create_non_recursive=0, gcs_api_client_gone_response_count=0, gs_filesystem_initialize=2, stream_read_vectored_incoming_ranges=0, stream_write_operations=0, gcs_list_dir_request=0, stream_read_operations=0, gcs_api_client_request_timeout_count=0, op_rename=0, op_get_file_status=1, op_glob_status=0, op_exists=0, stream_write_bytes=3209642, op_xattr_list=0, op_get_delegation_token=0, gcs_api_server_unavailable_count=0, directories_created=1, files_delete_rejected=0, stream_read_vectored_combined_ranges=0, op_xattr_get_named=0, gcs_list_file_request=2, op_hsync=0, action_http_get_request=0, stream_read_operations_incomplete=0, op_delete=0, stream_read_bytes=0, gcs_api_client_non_found_response_count=15, op_list_located_status=0, gcs_api_client_requested_range_not_satisfiable_count=0, op_hflush=17, op_list_status=0, stream_read_vectored_read_bytes_discarded=0, op_xattr_get_named_map=0, gcs_api_client_side_error_count=16, op_get_file_checksum=0, gcs_api_server_internal_error_count=0, stream_read_seek_bytes_skipped=0, stream_write_close_operations=0, gcs_get_media_request=0, gcs_connector_time=1862, files_deleted=0, action_http_post_request=7, op_mkdirs=1, gcs_api_client_rate_limit_error_count=0, op_copy_from_local_file=0, gcs_api_server_bad_gateway_count=0, stream_readVectored_range_duration=0, stream_read_seek_backward_operations=0, gcs_api_server_side_error_count=0, stream_read_seek_operations=0, gcs_get_other_request=2, stream_read_seek_forward_operations=0, gcs_api_client_precondition_failed_response_count=1, op_xattr_get_map=0, delegation_tokens_issued=0, gcs_backoff_time=0, gcs_list_dir_request_min=0, gcs_metadata_request_min=19, op_delete_min=0, op_glob_status_min=0, op_create_non_recursive_min=0, op_hsync_min=0, op_xattr_get_named_min=0, op_xattr_get_named_map_min=0, op_hflush_min=0, op_xattr_list_min=0, action_http_put_request_min=93, op_open_min=0, gcs_list_file_request_min=23, stream_write_close_operations_min=0, op_create_min=112, action_http_delete_request_min=30]
```

5 17°C Sunny

Search

04:15 PM 28-12-2025

Dashboard – qwiklabs-gcp-03-9... X +

https://console.cloud.google.com/home/dashboard?project=qwiklabs-gcp-03-9a342013a3da&pli=1&cloudshell=true

Free AI Paraphrasing... Transcript of Case St... GCP-LAB Hydra datawarehouse Products

Google Cloud qwiklabs-gcp-03-9a342013a3da Search (/) for resources, docs, products, and more Search

CLOUD SHELL Terminal (qwiklabs-gcp-03-9a342013a3da) X + Open Editor

Gemini CLI is available in Cloud Shell terminal! Type gemini to try it. [Learn more](#)

placement:  
clusterName: example-cluster  
clusterUuid: 3ef89636-2f20-4011-ac18-2c8259ea8458  
reference:  
jobId: 7230837507b64798a2feaa535de063e7  
projectId: qwiklabs-gcp-03-9a342013a3da  
sparkJob:  
args:  
- '1000'  
jarFileUris:  
- file:///usr/lib/spark/examples/jars/spark-examples.jar  
mainClass: org.apache.spark.examples.SparkPi  
status:  
state: DONE  
stateStartTime: '2025-12-28T10:45:32.456409Z'  
statusHistory:  
- state: PENDING  
stateStartTime: '2025-12-28T10:44:53.592055Z'  
- state: SETUP\_DONE  
stateStartTime: '2025-12-28T10:44:53.635936Z'  
- details: Agent reported job success  
state: RUNNING  
stateStartTime: '2025-12-28T10:44:53.956122Z'  
yarnApplications:  
- name: Spark Pi  
progress: 1.0  
state: FINISHED  
trackingUrl: http://example-cluster-m.us-east1-c.qwiklabs-gcp-03-9a342013a3da.internal.:8088/proxy/application\_1766918486251\_0001/  
student\_02\_77569a450c62@cloudshell:~ (qwiklabs-gcp-03-9a342013a3da)\$

5 17°C Sunny

Search

04:16 PM 28-12-2025

Dashboard – qwiklabs-gcp-03-9... X +

https://console.cloud.google.com/home/dashboard?project=qwiklabs-gcp-03-9a342013a3da&pli=1&cloudshell=true

Free AI Paraphrasin... Transcript of Case St... GCP-LAB Hydra datawarehouse Products

Google Cloud qwiklabs-gcp-03-9a342013a3da Search (/) for resources, docs, products, and more Search

CLOUD SHELL Terminal (qwiklabs-gcp-03-9a342013a3da) X + Open Editor

Gemini CLI is available in Cloud Shell terminal! Type gemini to try it. [Learn more](#)

placement:  
clusterName: example-cluster  
clusterUuid: 3ef89636-2f20-4011-ac18-2c8259ea8458  
reference:  
jobId: 7230837507b64798a2feaa535de063e7  
projectId: qwiklabs-gcp-03-9a342013a3da  
sparkJob:  
args:  
- '1000'  
jarFileUris:  
- file:///usr/lib/spark/examples/jars/spark-examples.jar  
mainClass: org.apache.spark.examples.SparkPi  
status:  
state: DONE  
stateStartTime: '2025-12-28T10:45:32.456409Z'  
statusHistory:  
- state: PENDING  
stateStartTime: '2025-12-28T10:44:53.592055Z'  
- state: SETUP\_DONE  
stateStartTime: '2025-12-28T10:44:53.635936Z'  
- details: Agent reported job success  
state: RUNNING  
stateStartTime: '2025-12-28T10:44:53.956122Z'  
yarnApplications:  
- name: Spark Pi  
progress: 1.0  
state: FINISHED  
trackingUrl: http://example-cluster-m.us-east1-c.qwiklabs-gcp-03-9a342013a3da.internal.:8088/proxy/application\_1766918486251\_0001/  
student\_02\_77569a450c62@cloudshell:~ (qwiklabs-gcp-03-9a342013a3da)\$

5 17°C Sunny

Search

04:16 PM 28-12-2025



00:26:12

Time limit

Caution: When you are in the console, do not deviate from the lab instructions. Doing so may cause your account to be blocked.

[Learn more](#)[Open Google Cloud console](#)

Username

student-02-77569a450c62



Password

1gd108qixhdc



Project ID

qwiklabs-gcp-03-9a34201



Student Resources

[Dataproc: Qwik Start - Qwiklabs](#)[Previous](#)

50/100

Lab instructions and task

GSP104

Overview

Setup and requirements

Task 1. Create a cluster

Task 2. Submit a job

Task 3. Update a cluster

Task 4. Test your understanding

Congratulations!

## Task 3. Update a cluster

- To change the number of workers in the cluster to four, run the following command:

```
gcloud dataproc clusters update example-cluster --num-workers 4
```

Your cluster's updated details are displayed in the command's output:

```
Waiting on operation [projects/qwiklabs-gcp-7f7aa0829e65200f/regions/global/operations/b86892cc-e71d-4e7b-aa5e-6030c945ea67].  
Waiting for cluster update operation...done.
```

- You can use the same command to decrease the number of worker nodes:

```
gcloud dataproc clusters update example-cluster --num-workers 2
```

Now you can create a Dataproc cluster and adjust the number of workers from the

[Next](#)

```
student_02_77569a450c62@cloudshell:~ (qwiklabs-gcp-03-9a342013a3da)$ gcloud dataproc clusters update example-cluster --num-workers 4
Waiting on operation [projects/qwiklabs-gcp-03-9a342013a3da/regions/us-east1/operations/5071845c-1138-3629-b647-04b9c1fade99].
Waiting for cluster update operation...done.
Updated [https://dataproc.googleapis.com/v1/projects/qwiklabs-gcp-03-9a342013a3da/regions/us-east1/clusters/example-cluster].
student 02 77569a450c62@cloudshell:~ (qwiklabs-gcp-03-9a342013a3da)$
```



00:23:32

Time limit

Caution: When you are in the console, do not deviate from the lab instructions. Doing so may cause your account to be blocked.  
[Learn more.](#)

[Open Google Cloud console](#)

Username

student-02-77569a450c62



Password

1gd108qixhdc



Project ID

qwiklabs-gcp-03-9a34201



Student Resources

[Dataproc: Qwik Start - Qwiklabs](#)

1. To change the number of workers in the cluster to four, run the following command:

```
gcloud dataproc clusters update example-cluster --num-workers 4
```

Your cluster's updated details are displayed in the command's output:

```
Waiting on operation [projects/qwiklabs-gcp-7f7aa0829e65200f/regions/global/operations/b86892cc-e71d-4e7b-aa5e-6030c945ea67].  
Waiting for cluster update operation...done.
```

2. You can use the same command to decrease the number of worker nodes:

```
gcloud dataproc clusters update example-cluster --num-workers 2
```

Now you can create a Dataproc cluster and adjust the number of workers from the `gcloud` command line on the Google Cloud.

Lab instructions and tasks

GSP104

Overview

Setup and requirements

Task 1. Create a cluster

Task 2. Submit a job

Task 3. Update a cluster

Task 4. Test your understanding

Congratulations!

[Previous](#)[Next](#)

Dashboard – qwiklabs-gcp-03-9... X +

https://console.cloud.google.com/home/dashboard?project=qwiklabs-gcp-03-9a342013a3da&pli=1&cloudshell=true

Free AI Paraphrasin... Transcript of Case St... GCP-LAB Hydra datawarehouse Products

Google Cloud qwiklabs-gcp-03-9a342013a3da Search (/) for resources, docs, products, and more Search

CLOUD SHELL Terminal (qwiklabs-gcp-03-9a342013a3da) X + Open Editor

Gemini CLI is available in Cloud Shell terminal! Type gemini to try it. [Learn more](#)

student\_02\_77569a450c62@cloudshell:~ (qwiklabs-gcp-03-9a342013a3da)\$ gcloud dataproc clusters update example-cluster --num-workers 4  
Waiting on operation [projects/qwiklabs-gcp-03-9a342013a3da/regions/us-east1/operations/5071845c-1138-3629-b647-04b9c1fade99].  
Waiting for cluster update operation...done.  
Updated [https://dataproc.googleapis.com/v1/projects/qwiklabs-gcp-03-9a342013a3da/regions/us-east1/clusters/example-cluster].  
student\_02\_77569a450c62@cloudshell:~ (qwiklabs-gcp-03-9a342013a3da)\$  
student\_02\_77569a450c62@cloudshell:~ (qwiklabs-gcp-03-9a342013a3da)\$ gcloud dataproc clusters update example-cluster --num-workers 2  
Waiting on operation [projects/qwiklabs-gcp-03-9a342013a3da/regions/us-east1/operations/db40ee2f-11ea-3b95-b364-d49258294110].  
Waiting for cluster update operation...done.  
Updated [https://dataproc.googleapis.com/v1/projects/qwiklabs-gcp-03-9a342013a3da/regions/us-east1/clusters/example-cluster].  
student\_02\_77569a450c62@cloudshell:~ (qwiklabs-gcp-03-9a342013a3da)\$

5 17°C Sunny

Search

04:22 PM 28-12-2025 ENG US